

mongoDB

Introducción a MongoDB

¡Bienvenidos!

Fernando Marín

fernando@mongodb.com


@fmarinperez




Agenda de hoy

- ¿Porqué existe NoSQL?
- Tipos de bases de datos NoSQL
- Características clave de MongoDB
- Tolerancia a fallos y persistencia de datos en MongoDB
- Escalabilidad en MongoDB
- Preguntas

El origen de SQL (70's)




El origen de SQL (70's)



250 Mb

\$ 81.000


El origen de SQL (70's)



250 Mb
\$ 81.000



El origen de SQL (70's)



250 Mb
\$ 81.000

Dennis Ritchie
Brian Kernighan
\$ 8.000/año ... (los dos)



Relational: RDBMS

PERSON

Pers_ID	Surname	First_Name	City
0	Miller	Paul	London
1	Ortega	Alvaro	Valencia
2	Huber	Urs	Zurich
3	Blanc	Gaston	Paris
4	Bertolini	Fabrizio	Rome

NO RELATION

CAR

Car_ID	Model	Year	Value	Pers_ID
101	Bently	1973	100000	0
102	Rolls Royce	1965	330000	0
103	Peugeot	1993	500	3
104	Ferrari	2005	150000	4
105	Renault	1998	2000	3
106	Renault	2001	7000	3
107	Smart	1999	2000	2

Relational: RDBMS



Relational: RDBMS



Relational: RDBMS → Row Stores



Relational: RDBMS → Row Stores



Relational: RDBMS → Row Stores



Relational: RDBMS → Row Stores



Relational: RDBMS → Row Stores





El origen de NoSQL (2000's)

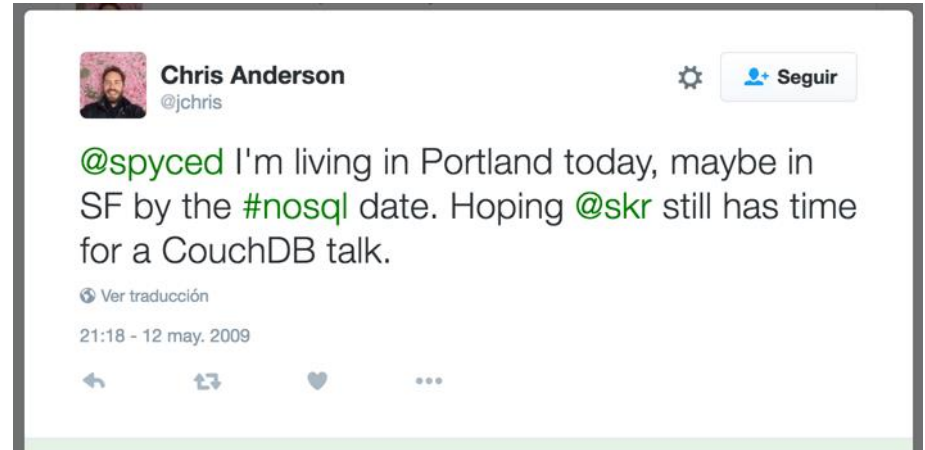
- Google y Amazon
 - Sistemas distribuidos
 - Open source
 - No relacionales

El origen de NoSQL (2000's)

- Google y Amazon
 - Sistemas distribuidos
 - Open source
 - No relacionales
- Mayo 2009 → twitter hashtag #nosql

El origen de NoSQL (2000's)

- Google y Amazon
 - Sistemas distribuidos
 - Open source
 - No relacionales



- Mayo 2009 → twitter hashtag #nosql



Tipos de Bases de Datos NoSQL

- Key-Value
- Wide Column
- Graph
- Document

Key Value Stores



Key Value Stores



Key Value Stores

- An associative array
- Single key lookup
- Very fast single key lookup

Key	Value
12345	4567.3456787
12346	{ addr1 : "The Grange", addr2: "Dublin" }
12347	"top secret password"
12358	"Shopping basket value : 24560"
12787	12345

Wide-Column Stores



Wide-Column Stores



Wide-Column Stores



Wide-Column Stores






Wide-Column Stores

- A series of consecutive seeks can retrieve a column efficiently
- Compressing similar data is super efficient
- So reads can grab more data off disk in a single seek
- How do I align my rows? By order or by inserting a row ID
- IF you just need a small number of columns you don't need to read all the rows
 - But → Updating and deleting by row is expensive
- Append only is preferred
- Better for OLAP than OLTP

Graph Stores

- Store graphs (edges and vertexes)
- E.g. social networks
- Designed to allow efficient traversal
- Optimised for representing connections
- Can be implemented as a key value stored with the ability to store links
- If your use case is not a graph you don't need a graph database

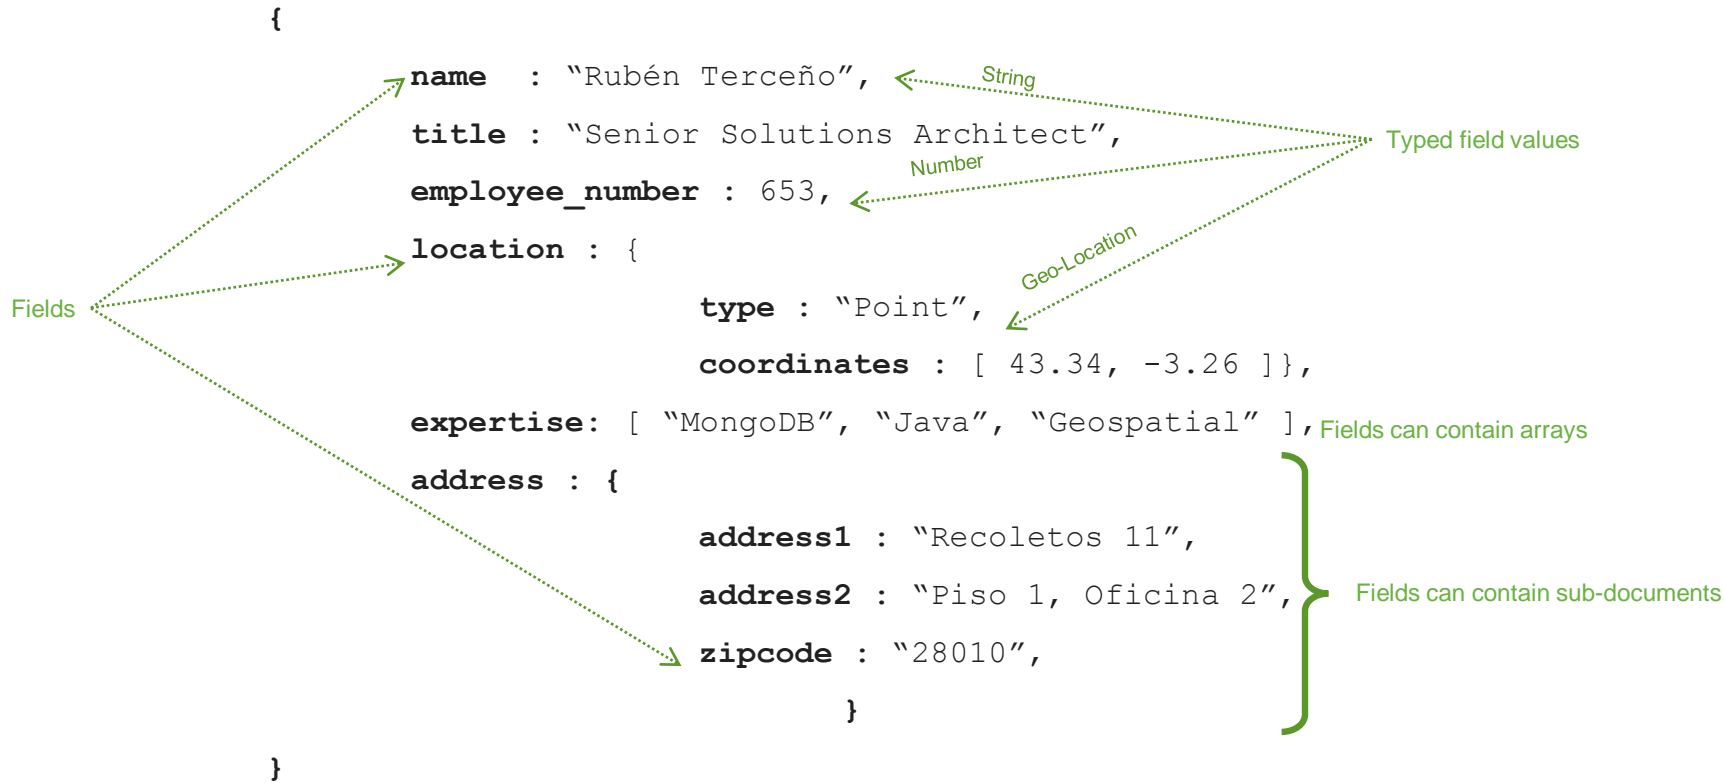


Document Store

- Not PDFs or Microsoft Word.
- Documents are rich structures like XML or Javascript Object Notation (JSON)

```
{  
  name : "Rubén Terceño",  
  title : "Senior Solutions Architect",  
  employee_number : 653,  
  location : {  
    type : "Point",  
    coordinates : [ 43.34, -3.26 ]},  
  expertise: [ "MongoDB", "Java", "Geospatial" ],  
  address : {  
    address1 : "Recoletos 11",  
    address2 : "Piso 1, Oficina 2",  
    zipcode : "28010",  
  }  
}
```

Documents are Rich Structures




Document Stores



Consistency vs. Availability




Consistency vs. Availability




Consistency vs. Availability




Consistency vs. Availability




Consistency vs. Availability




Consistency vs. Availability




Consistency vs. Availability




Consistency vs. Availability



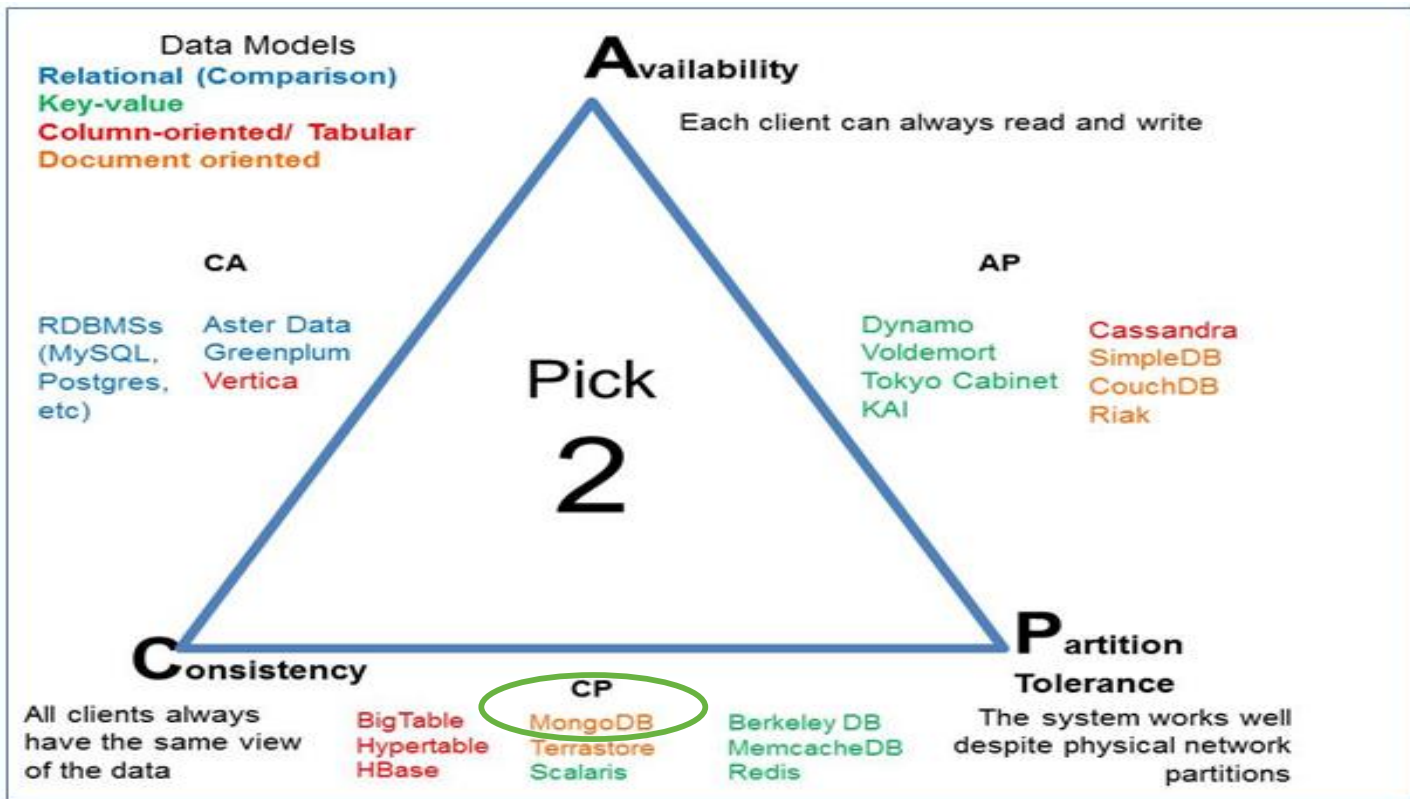
Consistency vs. Availability



Consistency vs. **Availability**



The CAP theorem





Agenda de hoy

- ~~¿Porqué existe NoSQL?~~
- ~~Tipos de bases de datos NoSQL~~
- Características clave de MongoDB
- Tolerancia a fallos y persistencia de datos en MongoDB
- Escalabilidad en MongoDB
- Preguntas

Relational




Relational

/ Expressive Query Language & Secondary Indexes

Expressive Query Language
& Secondary Indexes


Strong Consistency

Enterprise Management
& Integrations




- Queries to access data in sophisticated ways
- Indexes for efficient access to data for reads/writes
- Table stakes for a database

Relational / Strong Consistency



- Provide the most up-to-date copy of the data
- Complicated to build an application around eventually consistent model, even for the most sophisticated teams

Relational / Enterprise Management & Integrations



Database is just one piece of the puzzle and it needs to fit into the enterprise IT stack:

- Streamline ops according to best practices
- Visualize data / schema to make informed decisions and ensure data integrity
- Encrypt data & restrict access to database in accordance with existing security infrastructure
- Employ standard BI and visualization tools

The World Has Changed

Data

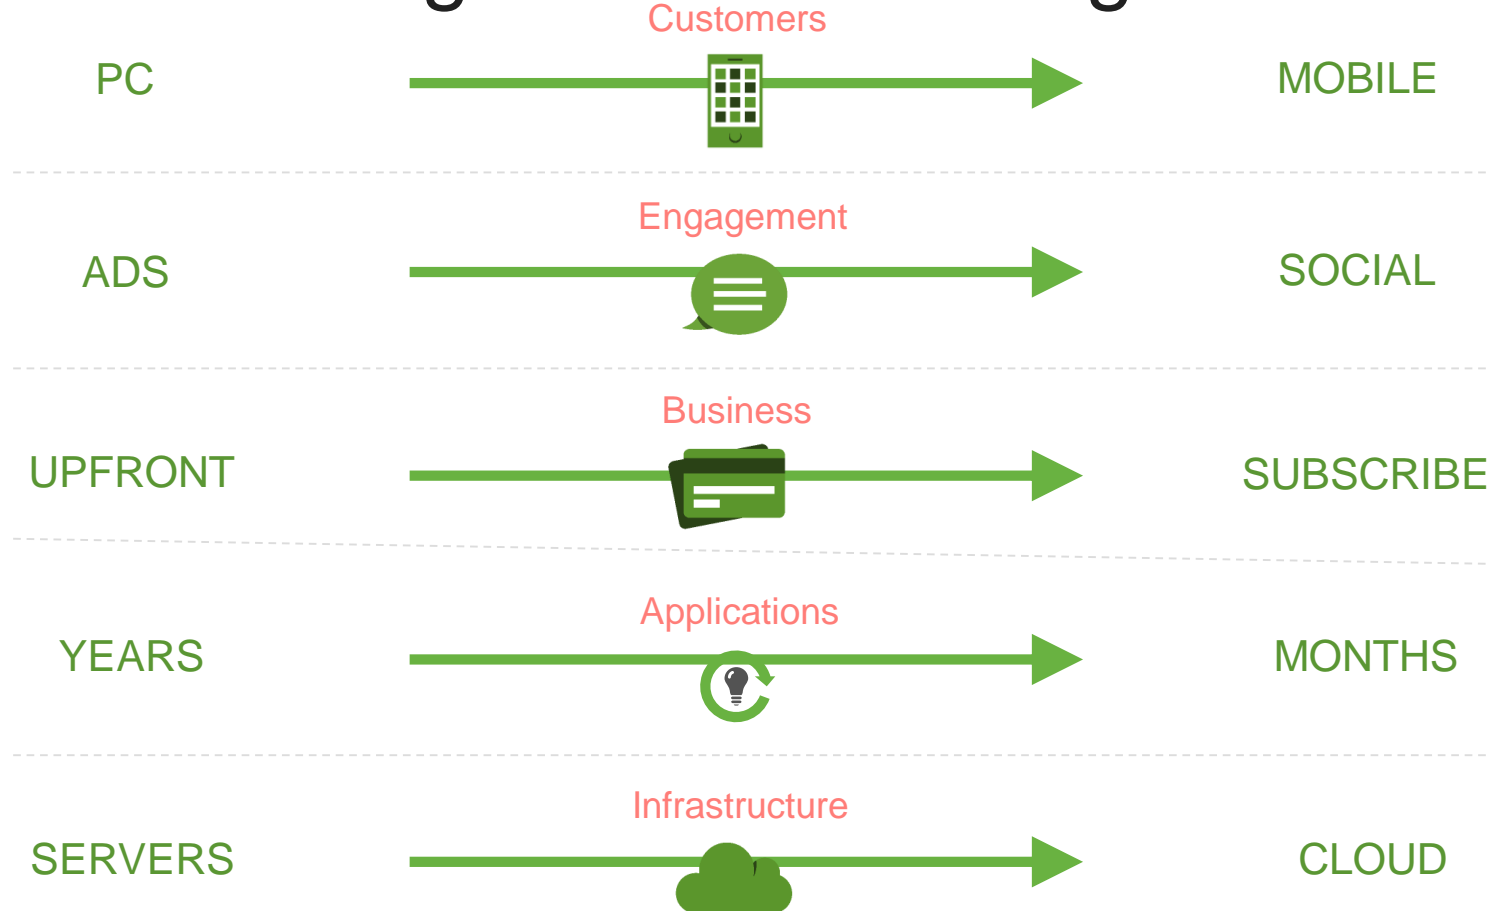
Risk

Time

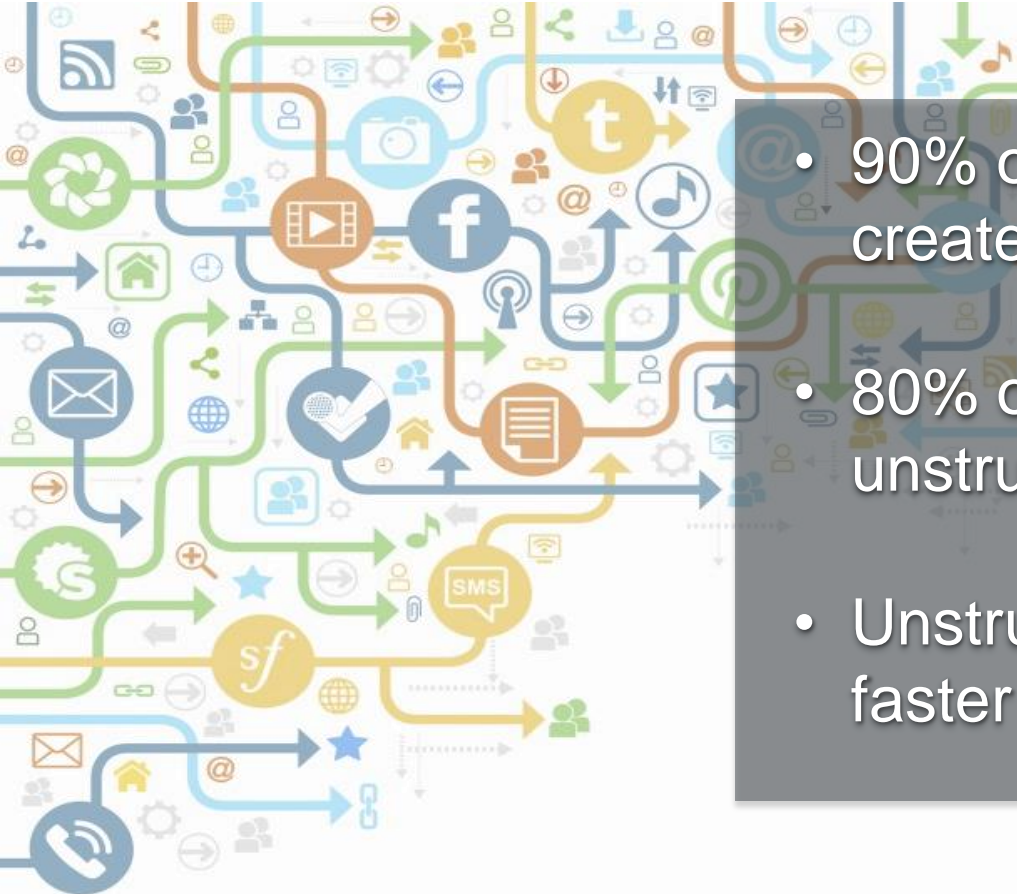
Cost



Information Management Has Changed




Your Data Has Changed



- 90% of the world's data was created in the last two years
- 80% of enterprise data is unstructured
- Unstructured data growing 2x faster than structured

Big Data Driving Factors

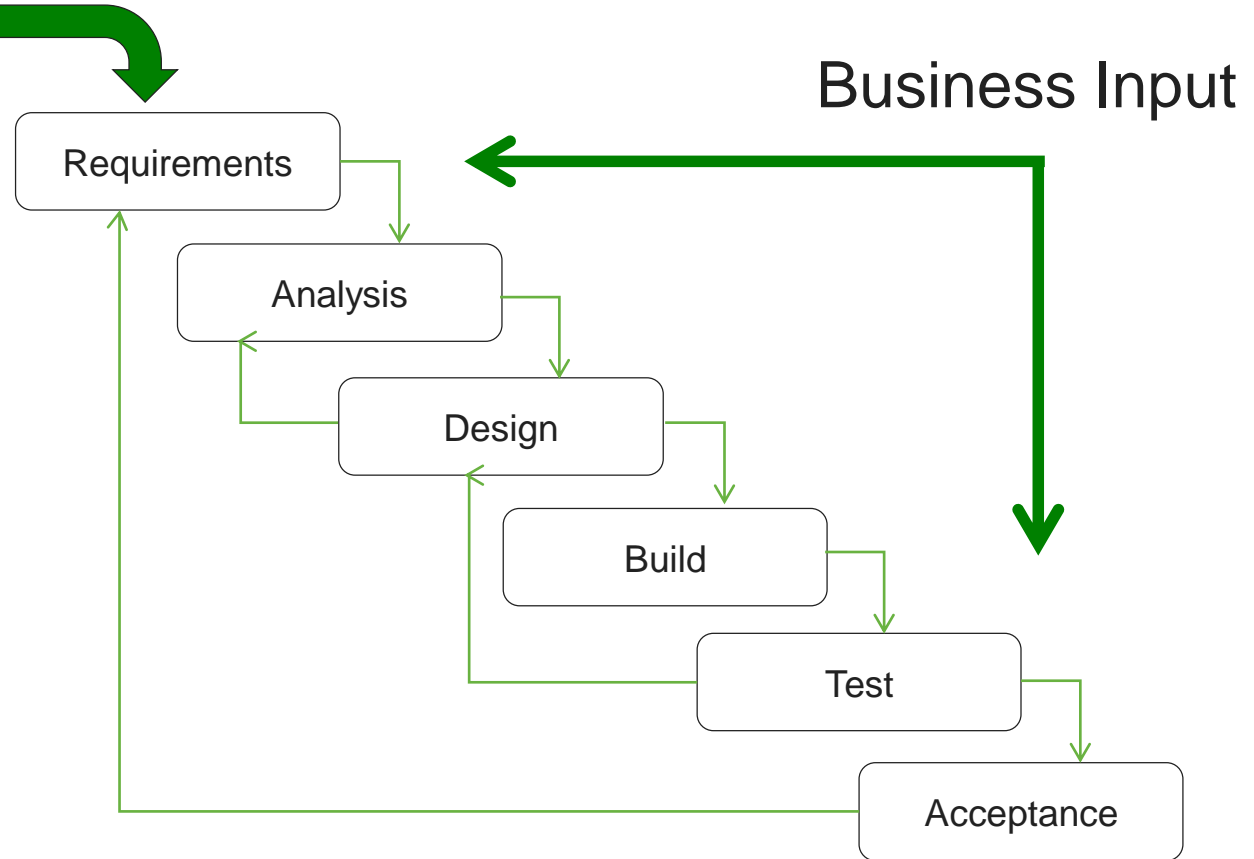
*What are the primary data issues driving you to consider Big Data?**



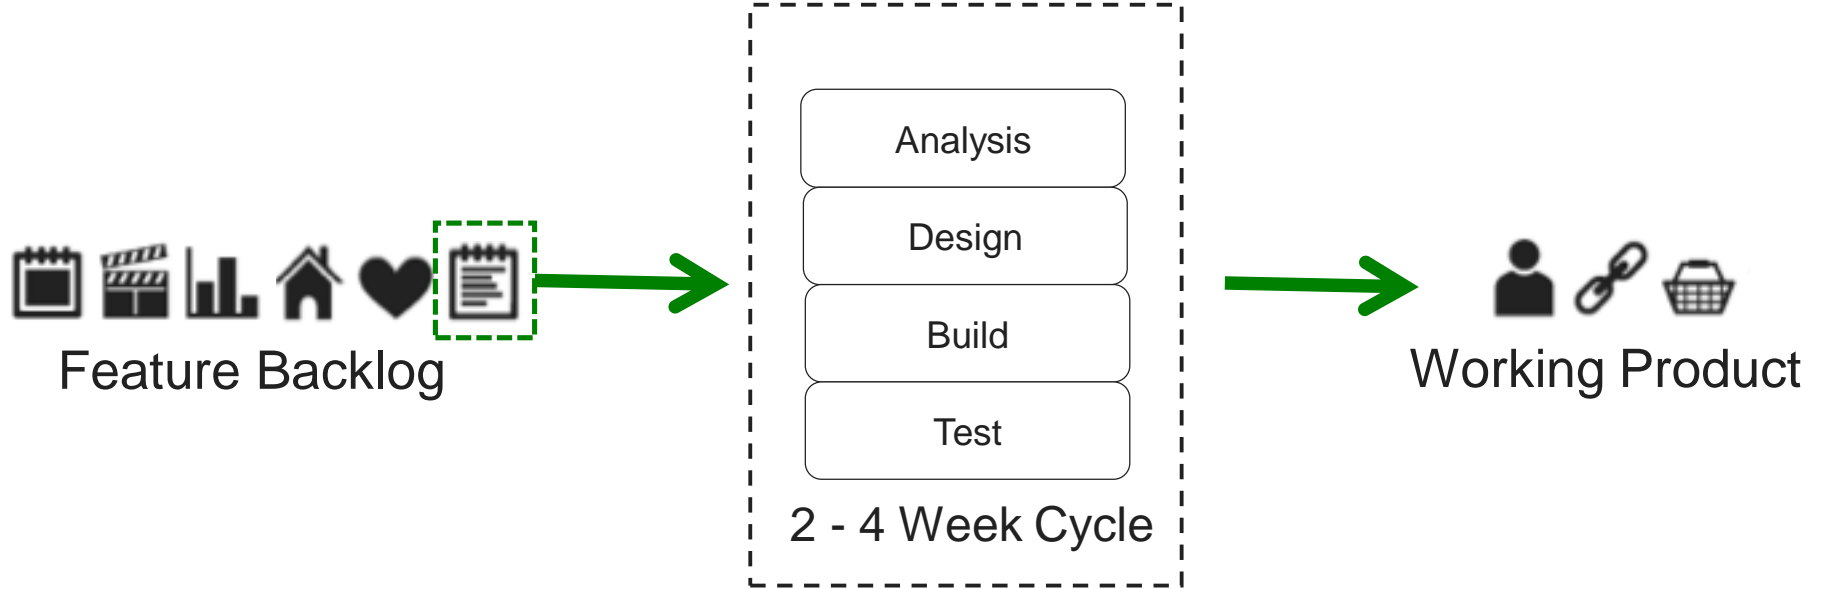
*“Of Gartner's "3Vs" of big data (volume, velocity, variety), the **variety of data sources** is seen by our clients as both the greatest challenge and **the greatest opportunity**.”*

FORRESTER

Development – Methods are Changing



Development – Agile Methodology



Software Has Changed



Microsoft

ORACLE®

- Higher up-front costs
- Higher TCO



mongoDB


redhat

- Lower up-front costs
- Lower TCO

NoSQL




NoSQL / Flexibility




- Allows faster iteration
- Accommodates various data types
- Data models include document, key-value, wide column

NoSQL / Scalability & Performance



- Scalability through partitioning / sharding
- NoSQL databases designed to deliver better performance than a relational database

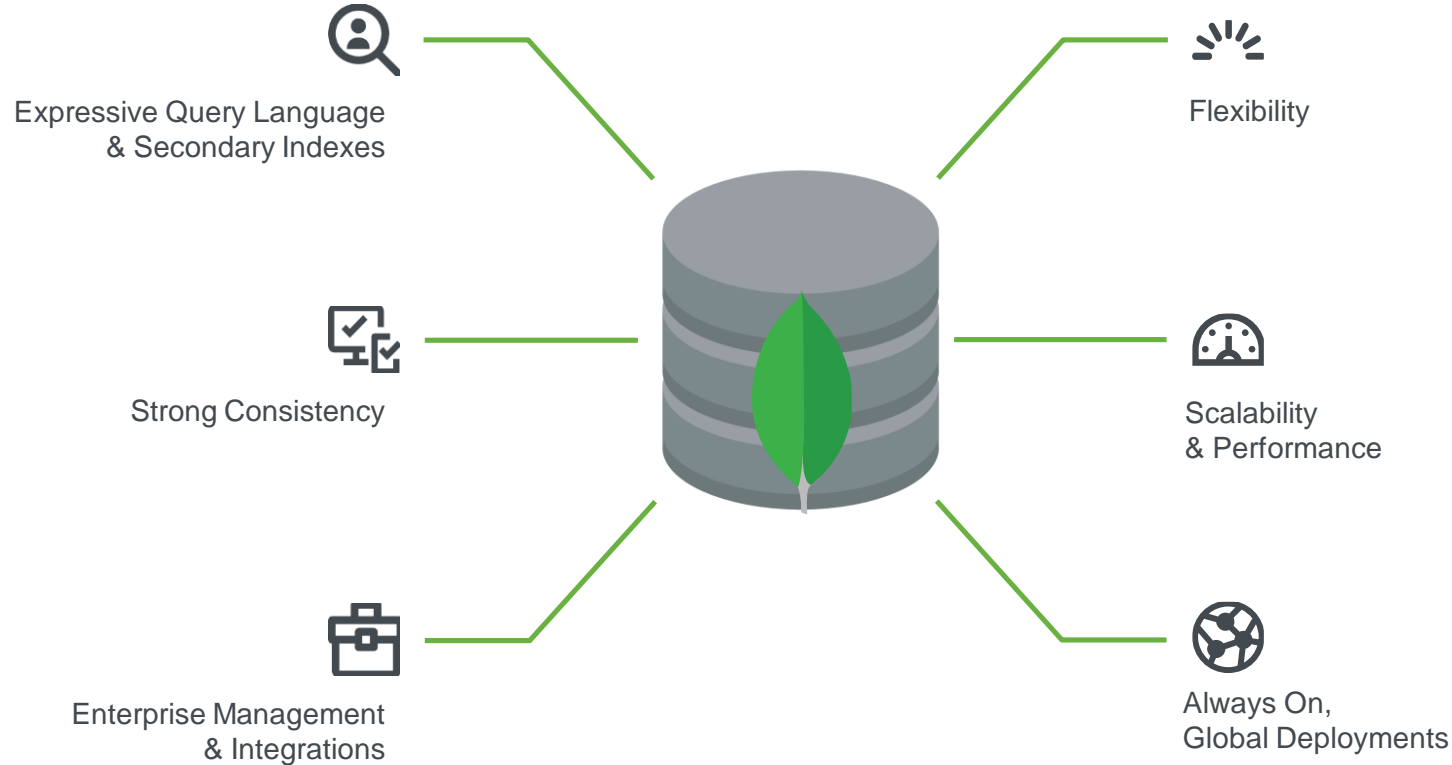
NoSQL / Always On, Global Deployments



Support highly available systems that provide a consistent, high-quality experience to end users around the world with:

- Run across many computers
- Automatically synchronize data across servers, racks, data centers

Nexus Architecture



Introducción a MongoDB



MongoDB


The leading NoSQL database



General
Purpose


```
{  
  name: "John Smith",  
  pfxs: ["Dr.", "Mr."],  
  address: "10 3rd St.",  
  phone: {  
    home: 1234567890,  
    mobile: 1234568138 }  
}
```

Document
Data Model



Open-
Source

MongoDB Business Value



Enabling New Apps



Faster Time to Market



Better Customer Experience



Lower TCO

Fortune 500 & Global 500

40 of the Top Financial Services Institutions

15 of the Top Retailers

15 of the Top Telcos

15 of the Top Technology Companies

15 of the Top Healthcare Companies

10 of the Top Electronics Companies

10 of the Top Media and Entertainment Companies

About MongoDB, Inc.



500+
employees



13 offices
worldwide



2,000+
customers



\$311M in
funding

MongoDB Use Cases

Single View



Internet of Things



Mobile



Real-Time Analytics



Catalog



Personalization



Content Management





4th Most Popular, Fastest Growing

RANK	DBMS	MODEL	SCORE	GROWTH (20 MO)
1.	Oracle	Relational DBMS	1,442	-5%
2.	MySQL	Relational DBMS	1,294	2%
3.	Microsoft SQL Server	Relational DBMS	1,131	-10%
4.	MongoDB	Document Store	277	172%
5.	PostgreSQL	Relational DBMS	273	40%
6.	DB2	Relational DBMS	201	11%
7.	Microsoft Access	Relational DBMS	146	-26%
8.	Cassandra	Wide Column	107	87%
9.	SQLite	Relational DBMS	105	19%

Source: [DB-engines database popularity rankings: May 2015](#)

The Largest Ecosystem

20,000,000+

MongoDB Downloads

400,000+

Online Education Students

40,000+

MongoDB Cloud Manager Users

35,000+

MongoDB User Group Members

1,000+


Technology and Services Partners

2,000+


Customers Across All Industries

Fastest-Growing Database

Google




LinkedIn




Twitter

Followers

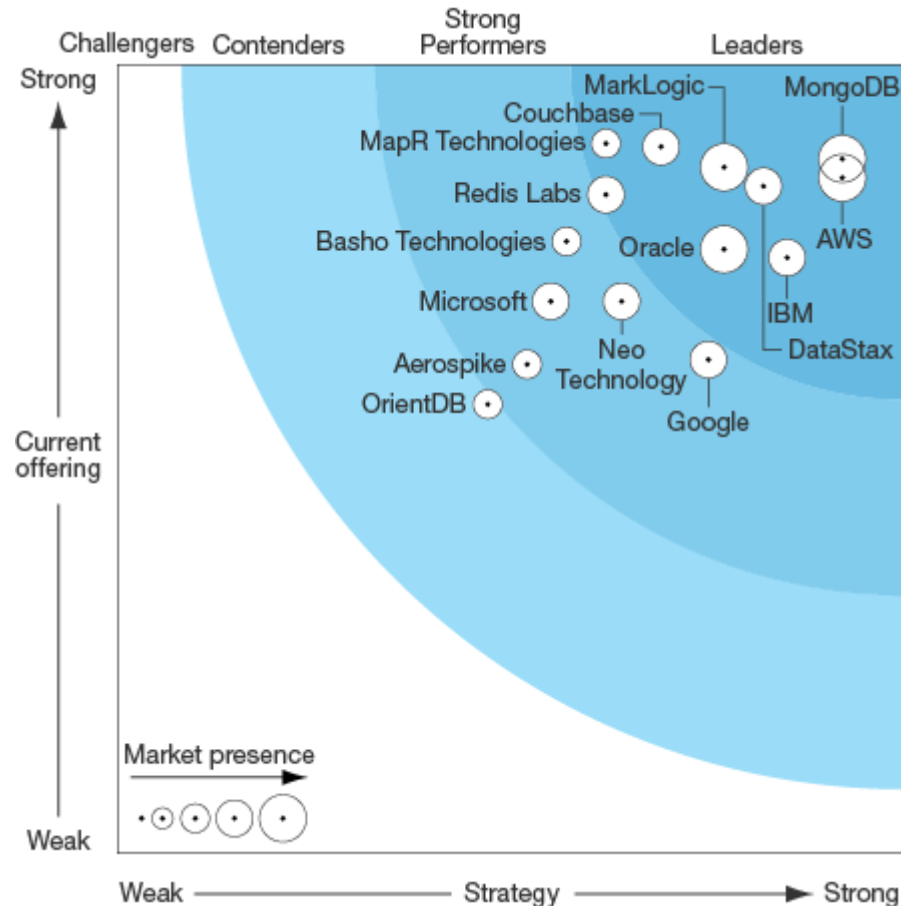


Facebook

Likes

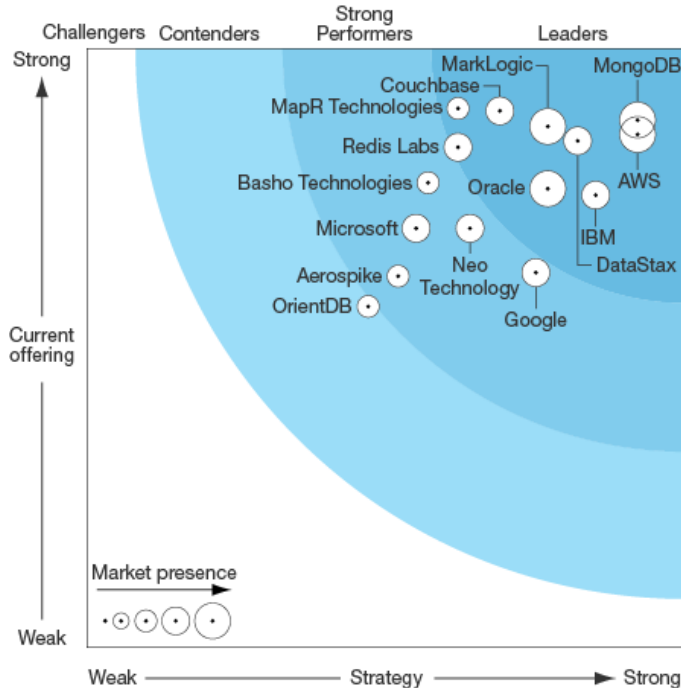


MongoDB named **A LEADER** in The Forrester Wave™: Big Data NoSQL, Q3 2016



The Forrester Wave™ is copyrighted by Forrester Research, Inc. Forrester and Forrester Wave are trademarks of Forrester Research, Inc. The Forrester Wave is a graphical representation of Forrester's call on a market and is plotted using a detailed spreadsheet with exposed scores, weightings, and comments. Forrester does not endorse any vendor, product, or service depicted in the Forrester Wave. Information is based on best available resources. Opinions reflect judgment at the time and are subject to change.

MongoDB named **a leader** in The Forrester Wave™: Big Data NoSQL, Q3 2016

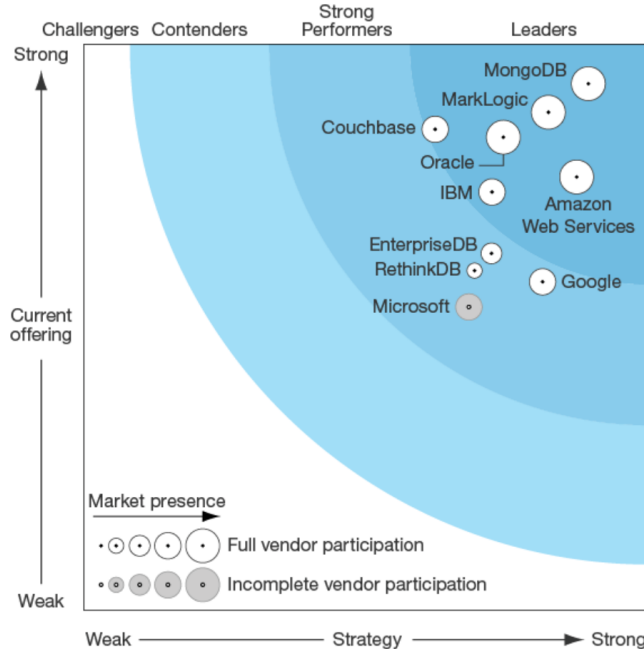


Achieving the highest possible scores in 19 of the 26 criteria, MongoDB believes that this report indicates its strength and versatility as an all-purpose database for a variety of mission-critical use cases.

5 out of 5 Points for:

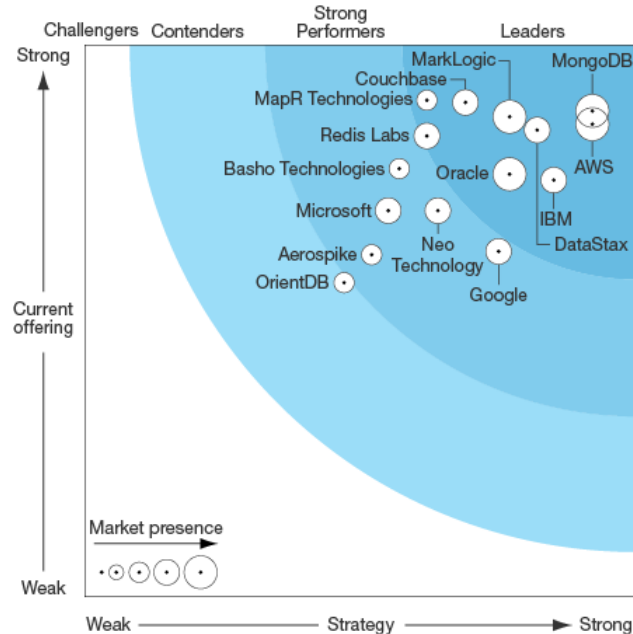
- Data Types
- Streaming and Loading
- Big Data Support
- In-memory
- Performance
- Scalability
- High Availability & Disaster Recovery
- Tools
- Workloads
- Use Cases
- Ability to Execute
- Road Map
- Open Source and Licensing
- Support

MongoDB named **a leader** in The Forrester Wave™: Document Stores, Q3 2016

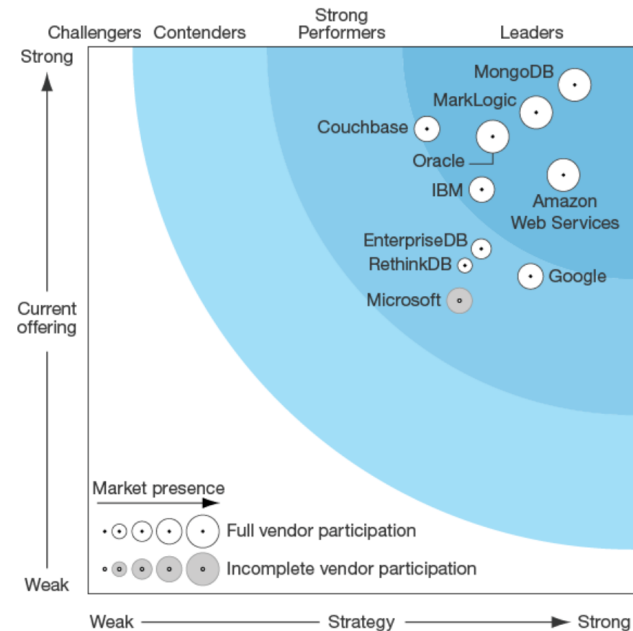


MongoDB received the highest scores possible in **20 of the 27** criteria evaluated, including performance, scalability, tooling, road map, support, product revenue and reach.

MongoDB named a **LEADER** in two key Forrester Waves:



The Forrester Wave™:
Big Data NoSQL, Q3 2016



The Forrester Wave™:
Document Stores, Q3 2016

Operational Database Landscape

Scalability & Performance


- Document Transactions
- Fast, Scalable Read/Writes

Depth of Functionality

Why other Databases Fall Short

Relational Model


Rigid Schemas
Resistant to
change



Data changes constantly,
In many use cases it's designed to

Scale-up

Throughput & Cost
make Scale-Up
Impractical




Scale-Up clusters were never meant
to handle today's volumes


How MongoDB Makes it Easy

Document Model

Scale-out



Flexible Multi-Structured Schema is
designed to adapt to changes



Scale-out to the end of The World and
distribute data where it needs to be

Relational Model

EmpID	Name	Dept	Title	Manager	Payband
9950	Dunham, Justin	Marketing	Product Manager	Neray, Graham	C

EmpBenPlanID	EmpFK	PlanFK
1	9950	100
2	9950	200

PlanID	Benefit	Plan
100	Health	PPO Plus
200	Dental	Standard

TitleID	Title
1500	Product Manager

BenID	Benefit
1	Health
2	Dental

DeptID	Department
500	Marketing

We Denormalise & Materialize for Performance & Simplicity

Relational Document Model

EmpID	Name	Dept	Title	Manager	Payband	
9950	Dunham, Justin	Marketing	Product Manager	Neray, Graham	Health	PPO Plus
					Dental	Standard



PlanID	Benefit
100	Health
200	Dental



```
{
  _id : 9950,
  employee_name: "Dunham, Justin",
  department : "Marketing",
  title : "Product Manager, Web",
  report_up: "Neray, Graham",
  pay_band: "C",
  benefits : [
    { type : "Health",
      plan : "PPO Plus" },
    { type : "Dental",
      plan : "Standard" }
  ]
}
```

We Denormalise & Materialize for Performance & Simplicity



No SQL But Still Flexible Querying

Rich Queries

- Find everybody who opened a special account last month in NY between \$100 and \$1000 OR last year more than \$500

Aggregation

- What is the average P&L of the trading desks grouped by a set of date ranges

Text Search

- Find all tweets that mention the bank within the last 2 days

Geospatial


- Find all customers that live within 10 miles of NYC

Map Reduce

- Calculate total amount settled position by symbol by settlement venue

How MongoDB Makes it Easy

In-Place Analytics



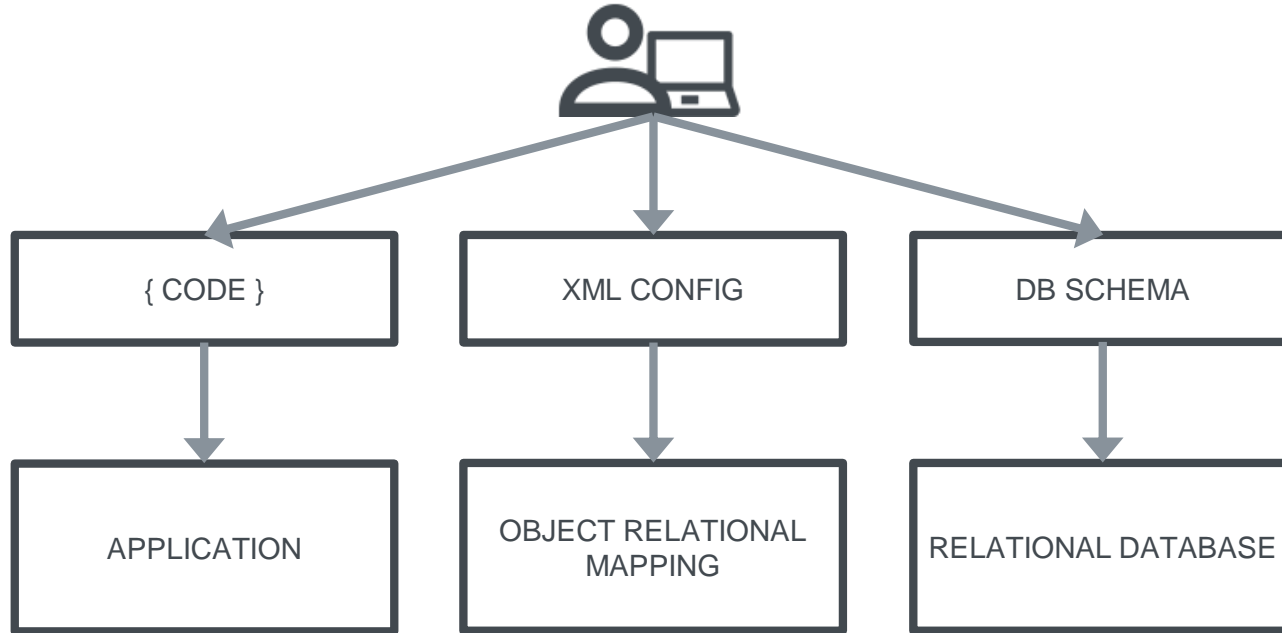
In Database Analytics provides application scalability

Security

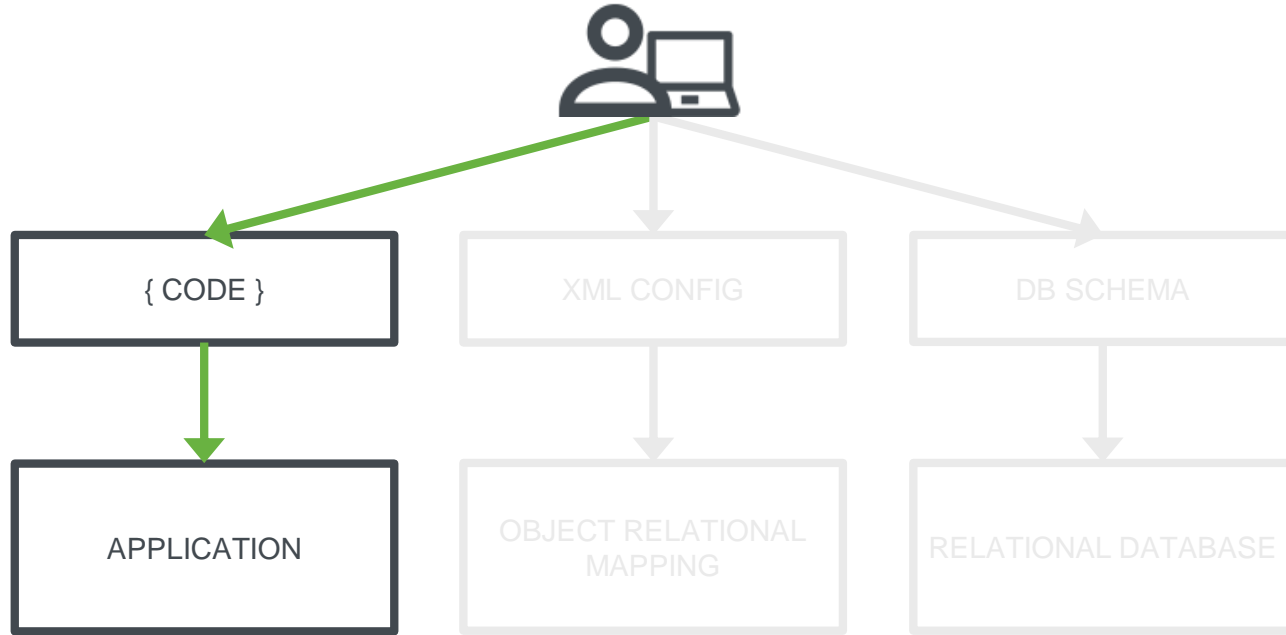


Authentication, Authorization, Auditing and Encryption protect your data

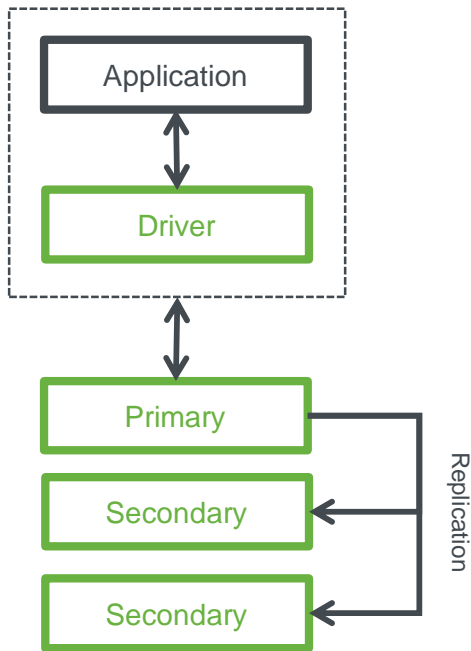
Development – The Past



Development – With MongoDB

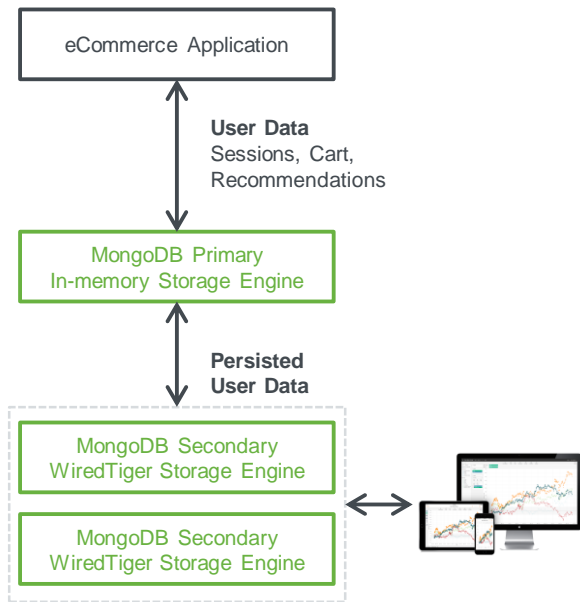


Replica Sets




- Replica set – 2 to 50 copies
- Replica sets make up a self-healing ‘shard’
- Data center awareness
- Replica sets address:
 - High availability
 - Data durability, consistency
 - Maintenance (e.g., HW swaps)
 - Disaster Recovery

Replica Sets – Workload Isolation

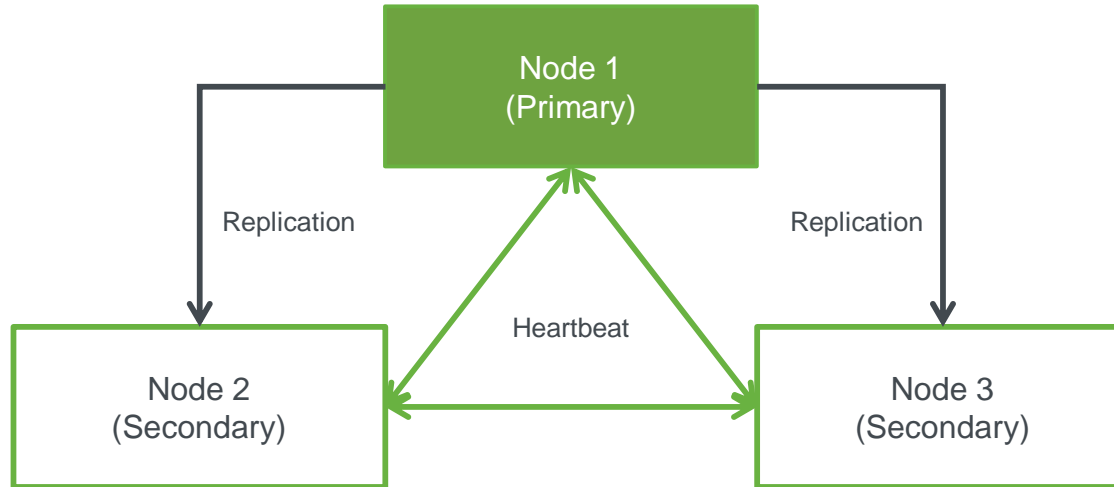


- Replica sets enable workload isolation
- Example: Operational workloads on the primary node, analytical workloads on the secondary nodes



Replica Set Creation





Replica Set - Initialize



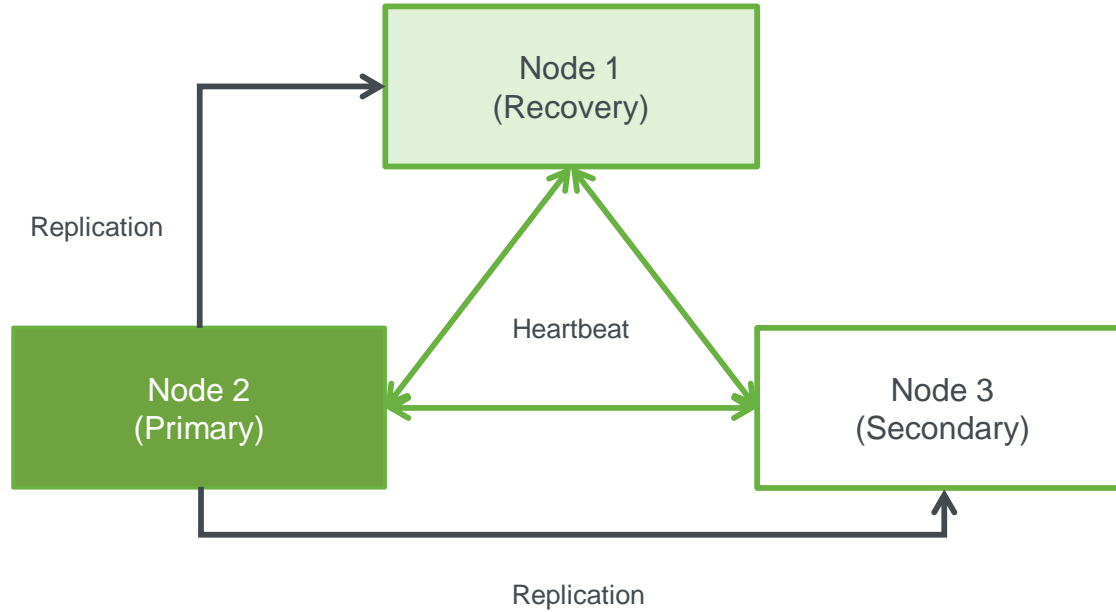
Replica Set - Failure



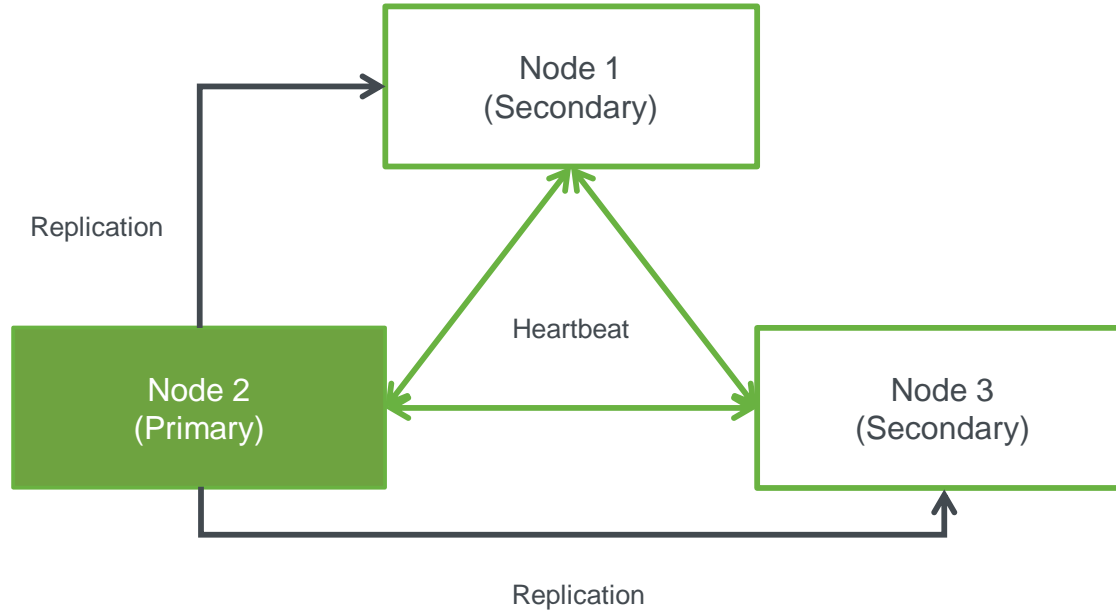
Replica Set - Failover




Replica Set - Recovery



Replica Set - Recovered



Elastic Scalability: Automatic Sharding



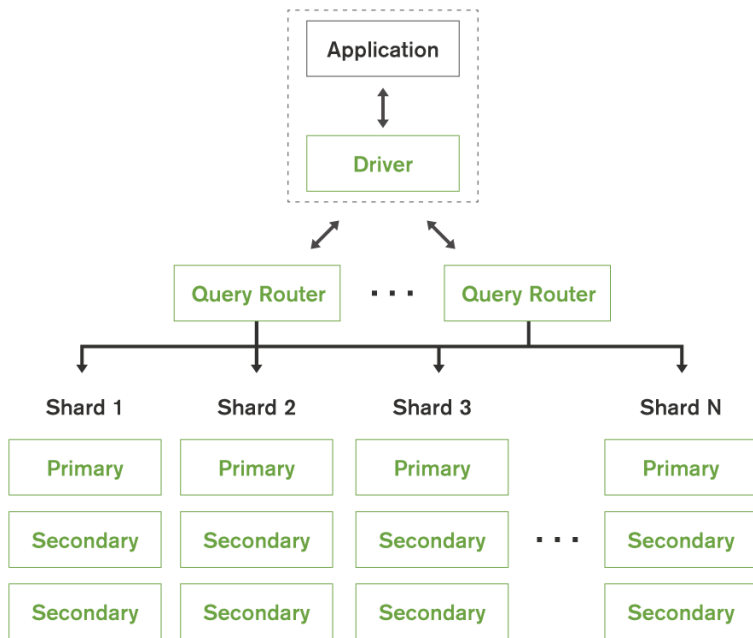
- **Increase or decrease** capacity as you go
- Automatic load balancing
- Three types of sharding
 - Hash-based
 - Range-based
 - Tag-aware



Scalability with Sharding

- Shard key partitions the content
- MongoDB automatically balances the cluster
- Shards can be added dynamically to a live system
- Rebalancing happens in the background
- Shard key is immutable
- Shard key can route queries to a specific shard
- Queries without a shard key are sent to all members
 - Each member process its part in parallel.

Query Routing



- Multiple query optimization models
- Each of the sharding options are appropriate for different apps / use cases



Query Routing

- With a sharded cluster we use a routing layer to guide queries
- We use a daemon called MongoS (Mongo Shard Router)
- Daemon is stateless
- Can run as many as required
- Typically one per app server



Agenda de hoy

- ~~¿Porqué existe NoSQL?~~
- ~~Tipos de bases de datos NoSQL~~
- ~~Características clave de MongoDB~~
- ~~Tolerancia a fallos y persistencia de datos en MongoDB~~
- ~~Escalabilidad en MongoDB~~
- Preguntas

¿Preguntas?

mongoDB